Dissertation Blueprint

Ma’at as Constitutional Infrastructure: Why All Advanced AI Systems Require a Ma’at-Governed Foundation

Suggested attribution page

• *Primary author / principal investigator:* Imhotep

• *Institutional development context:* Tehuti Research Lab

• *Research and drafting support:* African Scientific Chronology (custom GPT)

• *Prepared with AI-assisted research synthesis:* ChatGPT

Suggested acknowledgment language

This dissertation blueprint was conceptually directed by Imhotep and developed within the research environment of the Tehuti Research Lab. Research synthesis and drafting support were provided by African Scientific Chronology, a custom GPT designed for historical-scientific and socio-political analysis of Kemet, Africa, and systems thinking in modern technological design.

Suggested academic-use note

For formal university submission, keep Imhotep as the sole human author unless your institution permits another format, and disclose AI assistance in the acknowledgments or methods section according to institutional policy.

1. Working title

Ma’at as Constitutional Infrastructure: A Decolonial, Moral, and Systems-Theoretic Argument for Ma’at-Governed AI

Alternative title: From Alignment to Constitution: Why Ma’at Is Necessary in the Design, Governance, and Operation of AI Infrastructure

2. Abstract

This dissertation argues that advanced AI systems should not be governed merely by post hoc “safety layers” or narrow model-alignment techniques, but by a deeper constitutional order. I propose that Ma’at—understood as a framework of truth, order, balance, justice, reciprocity, and accountable action—offers the most coherent foundation for such an order. Ancient Egyptian/Kemetic traditions consistently associate Ma’at with truth, justice, balance, and cosmic-social order, while modern African-centered scholarship, especially Maulana Karenga’s work, frames Ma’at as a serious moral ideal capable of speaking to modern ethical discourse.

The dissertation situates this claim within the current AI governance crisis. NIST’s AI Risk Management Framework, its Generative AI Profile, UNESCO’s Recommendation on the Ethics of Artificial Intelligence, the OECD AI Principles, and OWASP’s LLM risk taxonomy all converge on the need for governance, transparency, human oversight, accountability, lifecycle monitoring, and risk management. Frontier developers are moving in the same direction: Anthropic now publishes a formal model constitution and frames frontier model capability as a cybersecurity issue tied to critical software defense. These developments show that the field increasingly recognizes AI as a problem of constitutional order, not just model capability.

My central claim is that Ma’at is needed because modern AI systems exercise organized power. They generate outputs, shape knowledge, influence decisions, actuate tools, and increasingly mediate institutional processes. Power of that kind cannot be left to raw intelligence, scale, or corporate preference alone. It must be bounded by a governing order that makes truth traceable, action reviewable, memory accountable, identity explicit, and authority distributed. Ma’at provides a conceptual and architectural grammar for that task.

3. Core thesis statement

This dissertation contends that Ma’at should function as the constitutional foundation of AI infrastructure because advanced AI systems are not merely predictive engines but sociotechnical regimes of organized power, and such regimes require a governing order grounded in truth, balance, justice, accountability, and disciplined self-correction.

Stronger version: Without a Ma'at-like constitutional layer, AI systems predictably drift toward epistemic distortion, operational opacity, unsafe autonomy, and moral incoherence; therefore, trustworthy AI requires Ma'at not as symbolism, but as infrastructure doctrine.

4. Research questions

Primary research question

How can Ma'at be translated from a classical Kemetic moral-constitutional ideal into a modern framework for governing AI infrastructure?

Secondary questions

• Why are current AI safety and ethics frameworks necessary but insufficient without a deeper constitutional theory?

• In what sense does Ma'at differ from generic "ethics layers" or post hoc alignment patches?

• How can Ma'at be operationalized across data, models, memory, orchestration, retrieval, tooling, and audit layers?

• What does a Ma'at-governed AI stack look like in practice?

• How does an African-centered governance model challenge the civilizational narrowness of current AI ethics discourse?

5. Main propositions

Proposition 1

AI governance has already moved beyond a purely technical paradigm and now operates as a constitutional problem involving transparency, accountability, oversight, and bounded authority. NIST, UNESCO, OECD, and OWASP all show this shift, even if they do not use that language.

Proposition 2

Ma'at is not merely a religious or symbolic concept; it is a moral-political ideal of truth, balance, justice, and order that can be translated into modern institutional design.

Proposition 3

A decolonial AI future requires drawing seriously on African intellectual traditions, not only on Euro-American regulatory categories. UNESCO itself emphasizes the importance of recognizing endogenous cultures, values, and knowledge in AI governance.

Proposition 4

Ma'at is best operationalized not at the model layer alone but across the full infrastructure stack: data provenance, identity, memory, event systems, policy gates, session awareness, audit trails, and human review.

6. Significance of the study

This dissertation makes five contributions.

1. *Decolonial intervention:* It offers a decolonial intervention into AI governance by taking an African civilizational framework seriously as a source of modern systems design rather than merely as heritage or symbolism. Diop and Obenga are relevant here because both worked to restore ancient Egypt/Kemet to African intellectual history and to defend African-centered historical and philosophical inquiry against exclusionary traditions.

2. *Organized power:* It reframes AI from "smart software" to organized power, which means that governance must be infrastructural, not ornamental.

3. *Constitutional model:* It contributes a constitutional model of AI, arguing that trustworthy systems require more than benchmark performance and post hoc moderation.

4. *Technical translation:* It provides a path for translating Ma’at into modern technical architecture: provenance, policy enforcement, anomaly tagging, role separation, and bounded action.

5. *Practical foundation:* It provides a foundation for practical implementations in environments such as the Tehuti Research Lab, where RAG, swarm systems, policy engines, and memory systems can be rebuilt on explicit moral-constitutional grounds.

7. Literature review

A. Ma’at as moral and constitutional order

The concept of Ma’at is consistently described in major reference sources as involving truth, balance, justice, and order. The British Museum characterizes Ma’at as the deified personification of order, balance, truth, and justice, and notes that maintaining Ma’at was tied to rightful rule and the maintenance of order. Britannica similarly summarizes Ma’at as truth, justice, and cosmic order.

For a modern African-centered philosophical treatment, Maulana Karenga’s Maat, The Moral Ideal in Ancient Egypt is essential. Its own publisher description makes clear that the book aims to present Ma’at in the language of modern moral discourse while preserving its distinctiveness as an ethical tradition capable of sustaining critical philosophical reflection. That makes Karenga especially important for any dissertation trying to move from ancient concept to modern institutional design.

Key claim for literature review: Ma’at already contains the conceptual elements that modern AI governance struggles to reunify: truthful order, proportion, right relation, duty, and accountable judgment.

B. African-centered restoration of Kemet as an intellectual source

Cheikh Anta Diop’s work is vital because it helped restore ancient Egypt to African history and challenged Eurocentric accounts of civilization and knowledge. Britannica summarizes Diop as a major scholar associated with arguments for the African nature of Egyptian civilization and the cultural unity of Africa.

Théophile Obenga deepened this restoration by combining historical, linguistic, and Egyptological scholarship. San Francisco State University identifies him as an Egyptologist, linguist, and historian who contributed to UNESCO’s General History of Africa; Présence Africaine describes him as working with Diop to recenter African history around African intellectual concerns.

Key claim for literature review: The dissertation stands within a stream of scholarship that refuses to treat Kemet as detachable from African thought, and therefore refuses to treat Ma’at as unavailable for modern theoretical work.

C. Contemporary AI governance frameworks

NIST’s AI RMF 1.0 organizes AI risk management around four functions—Govern, Map, Measure, and Manage—and explicitly treats governance as cross-cutting. NIST’s Generative AI Profile extends that logic to generative systems and aims to help organizations incorporate trustworthiness into design, development, use, and evaluation. NIST has also now released a concept note for a critical-infrastructure AI profile, showing that AI governance is increasingly being treated as infrastructure governance.

UNESCO’s Recommendation frames human rights, dignity, transparency, fairness, and human oversight as foundational, and explicitly states that monitoring and assessment should be continuous and proportionate to risk. UNESCO also emphasizes the recognition and promotion of endogenous cultures, values, and knowledge, which is especially relevant for a Ma’at-centered dissertation.

The OECD AI Principles likewise foreground human rights, fairness, transparency, explainability, robustness, security, safety, and accountability.

Key claim for literature review: Mainstream governance frameworks already identify the right risk categories, but they usually remain fragmented. Ma’at contributes the deeper unity that these frameworks imply but do not fully theorize.

D. AI security, agentic risk, and the constitutional turn

OWASP’s Top 10 for LLM Applications shows that modern AI risk is no longer limited to hallucination. Prompt injection, insecure output handling, training data poisoning, insecure plugin design, excessive agency, and overreliance are now standard categories of AI-system failure.

Anthropic's public materials show that frontier labs are increasingly adopting constitutional language themselves. Anthropic's 2026 constitution is described as a foundational document shaping model behavior and treated as the final authority governing other instructions. Anthropic also now frames frontier-model cyber capability as a watershed security issue through Project Glasswing and Mythos Preview.

Key claim for literature review: The field is already moving toward constitutional governance. The question is not whether AI needs a constitution. The question is which constitution, grounded in what moral order.

8. Theoretical framework

This dissertation proposes *Ma'at as constitutional infrastructure*.

That phrase means Ma'at is not treated as a decorative ethical theme placed on top of AI. It is treated as the governing logic that should shape system boundaries, identities, permissions, memory, claims to truth, and routes of action.

Ma'at is defined operationally through six infrastructural principles:

1. *Truth — No output, record, or action should be treated as authoritative without provenance, traceability, and reviewability. 2. Balance — Confidence, autonomy, and privilege should be proportional to evidence, context, and risk. 3. Order — Systems should preserve clear contracts: identity, event schema, memory schema, policy schema, and role boundaries. 4. Justice — Access, action, and enforcement should be accountable, non-arbitrary, and reviewable. 5. Reciprocity — Systems should remain in right relation to users, institutions, and affected communities, rather than treating optimization as its only goal. 6. Accountability* — No meaningful action should be invisible. Actions must be attributable to actors, sessions, policies, and records.

This framework translates Ma'at into modern AI architecture without collapsing it into generic compliance language.

9. Methodology

This project is best described as a *normative-architectural dissertation* grounded in comparative ethics, decolonial theory, and systems design.

Methodological components

A. Conceptual reconstruction — Reconstruct Ma'at from major historical and philosophical treatments, especially classical descriptions and African-centered ethical scholarship.

B. Comparative governance analysis — Compare Ma'at's moral-constitutional structure to NIST, UNESCO, OECD, and OWASP governance concerns.

C. Infrastructure translation — Translate Ma'at into architectural requirements across: data, model, retrieval, memory, orchestration, tooling, audit, incident response.

D. Case-based systems reflection — Use contemporary agentic and frontier-AI developments as evidence that unguided intelligence creates governance crises. Anthropic's public constitutional and security materials are especially useful here because they show the industry already recognizes the need for explicit governing documents and safeguards.

Methods statement

The dissertation is not an empirical benchmark paper. It is a theoretical and design-oriented dissertation that develops a constitutional model for AI and tests its coherence against existing governance, security, and institutional frameworks.

Limitations

This project does not claim that Ma'at is the only possible ethical foundation for AI. It claims that a Ma'at-like constitutional order is necessary and that Ma'at itself provides one of the strongest available models for such order.

10. Chapter-by-chapter outline

Chapter 1 — The Crisis of Ungoverned Intelligence

Chapter thesis: AI has outgrown the assumption that intelligence alone is enough; advanced systems now require constitutional order.

Sections: The shift from software to infrastructure; Why “alignment” is too narrow a category; Frontier capability and civilizational risk; The limits of patchwork safety.

Core sources: NIST AI RMF, NIST GAI Profile, OWASP, Anthropic constitutional and cyber materials.

Chapter 2 — Ma’at: Truth, Balance, Justice, and Ordered Power

Chapter thesis: Ma’at should be understood as a constitutional moral order, not merely an ancient religious symbol.

Sections: Ma’at in classical description; Ma’at as ethical ideal; Ma’at and rightful governance; Why Ma’at is structurally relevant to AI.

Core sources: British Museum, Britannica, Karenga.

Chapter 3 — Kemet, Africa, and the Decolonization of AI Ethics

Chapter thesis: AI ethics remains civilizationally narrow unless African frameworks like Ma’at are treated as sources of theory, not as optional cultural supplements.

Sections: Eurocentrism in the genealogy of AI ethics; Diop and the restoration of Kemet to African history; Obenga and African intellectual continuity; UNESCO and endogenous knowledge systems.

Core sources: Diop, Obenga, UNESCO.

Chapter 4 — Why Existing AI Governance Is Necessary but Incomplete

Chapter thesis: Contemporary frameworks correctly identify many risks but lack a unified constitutional grammar.

Sections: NIST and lifecycle risk; UNESCO and human responsibility; OECD and trustworthy AI; OWASP and applied system risk; The fragmentation problem.

Core sources: NIST, UNESCO, OECD, OWASP.

Chapter 5 — Ma’at as Constitutional Infrastructure

Chapter thesis: Ma’at provides a framework for governing data, memory, action, and accountability across the AI stack.

Sections: Truth as provenance; Balance as proportional autonomy; Order as schema and role discipline; Justice as bounded authority; Reciprocity as human-centered relation; Accountability as event visibility and auditability.

Contribution: This is the theory chapter where the dissertation’s original framework is formally stated.

Chapter 6 — Building Ma’at into AI Systems

Chapter thesis: Ma’at can be operationalized through concrete infrastructure patterns.

Sections: Identity contracts; Event contracts; Structured memory and provenance; Retrieval with contradiction tracking; Policy engines and Guard systems; Session indexing and live situational awareness; Anomaly tagging and quarantine logic.

Contribution: This chapter bridges philosophy and engineering.

Chapter 7 — From Alignment to Constitution: Tehuti Guard and the Ma’at Stack

Chapter thesis: Model behavior should be governed by a constitutional control plane rather than trusted as self-governing intelligence.

Sections: The model as proposer, not sovereign; Policy gating before execution; Human review and bounded agency; Tehuti Guard as constitutional enforcement; Swarm roles and distributed responsibility; Why no single component should own truth.

Contribution: This chapter translates the dissertation into a working architecture for labs, enterprises, and civic systems.

Chapter 8 — Toward the 42 Laws of Ma’at for Technology

Chapter thesis: The future of AI governance requires not just principles but codified operating laws.

Sections: Why modern AI needs explicit moral-operational laws; Drafting a Ma’at code for technology; Principles for education, research, and SaaS deployment; Future research agenda.

Appendix tie-in: This chapter can lead directly into a formal appendix: 42 Laws of Ma’at for AI and Technology.

Chapter 9 — Conclusion: The Future of Trustworthy AI Is Constitutional

Chapter thesis: The decisive question is no longer how powerful AI can become, but what governing order will contain and direct that power.

Sections: Summary of argument; Contributions to AI ethics and governance; Contributions to African philosophy and STS; Contributions to AI infrastructure design; Final claim: Ma’at as the missing constitutional grammar of the AI age.

11. Expected contributions to the field

Domain	Contribution
Philosophical	A new account of Ma’at as a modern constitutional resource for technological systems.
Governance	A unified framework that links provenance, accountability, oversight, and risk into one moral-operational order.
Decolonial	A challenge to the civilizational narrowness of prevailing AI ethics.
Engineering	A translation of moral theory into infrastructure components: identity, policy, event logging, memory structure, anomaly tagging, and bounded execution.
Institutional	A framework usable in labs, education systems, businesses, and public institutions.

12. Proposed appendices

- *Appendix A:* Draft 42 Laws of Ma’at for Technology
- *Appendix B:* Identity schema for governed AI systems
- *Appendix C:* Event schema and audit taxonomy
- *Appendix D:* Memory provenance schema for RAG systems
- *Appendix E:* Tehuti Guard control-flow model
- *Appendix F:* Session Index specification
- *Appendix G:* Research ethics and AI-assistance disclosure statement

13. Selected bibliography

Foundational Ma’at and African-centered sources

- British Museum. Maat object entry.
- Britannica. Maat entry.
- Karenga, Maulana. Maat: The Moral Ideal in Ancient Egypt: A Study in Classical African Ethics.
- Diop, Cheikh Anta. Selected works on African history, civilization, and Egypt.
- Obenga, Théophile. Works on African philosophy, linguistics, and Egyptology.

AI governance and ethics

- NIST. AI Risk Management Framework 1.0.
- NIST. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).
- UNESCO. Recommendation on the Ethics of Artificial Intelligence.
- OECD. AI Principles.

AI security and agentic risk

- OWASP. Top 10 for Large Language Model Applications.
- Anthropic. Claude’s new constitution.
- Anthropic Frontier Red Team. Assessing Claude Mythos Preview’s cybersecurity capabilities.

Internal corpus to cite directly in the final dissertation

If you intend the final dissertation to reflect the specific Tehuti Research Lab and UKMT Press line of argument, cite those texts directly in the finished version rather than relying on summaries:

- What Must Be Done
- Restoring History
- BAK2
- African Time

These texts were not cited in detail here because their contents were not provided in full.

14. Final framing paragraph (reusable)

This dissertation argues that the age of AI has become an age of constitutional design. Advanced systems now mediate knowledge, action, memory, and institutional power. As such, they cannot be governed by capability metrics and patchwork safety measures alone. They require a deeper order—one that binds truth to evidence, action to accountability, autonomy to proportional restraint, and intelligence to moral legitimacy. Ma’at provides precisely that order. It should therefore be treated not as an optional ethical ornament, but as constitutional infrastructure for the AI age.

Optional next step: develop Chapter 1 in full dissertation prose with citations and footnote-ready formatting.